

## 3.RF-Grid

October 13, 2022

```
[1]: #importing the Libraies
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
[2]: # Reading the Dataset
dataset = pd.read_csv('insurance_pre.csv')
```

```
[ ]: dataset
```

```
[ ]: dataset
```

```
[3]: dataset=pd.get_dummies(dataset,drop_first=True)
```

```
[4]: dataset
```

```
[4]:
```

	age	bmi	children	charges	sex_male	smoker_yes
0	19	27.900	0	16884.92400	0	1
1	18	33.770	1	1725.55230	1	0
2	28	33.000	3	4449.46200	1	0
3	33	22.705	0	21984.47061	1	0
4	32	28.880	0	3866.85520	1	0
...	...	...	...	...	...	...
1333	50	30.970	3	10600.54830	1	0
1334	18	31.920	0	2205.98080	0	0
1335	18	36.850	0	1629.83350	0	0
1336	21	25.800	0	2007.94500	0	0
1337	61	29.070	0	29141.36030	0	1

[1338 rows x 6 columns]

```
[5]: indep=dataset[['age', 'bmi', 'children', 'sex_male', 'smoker_yes']]
dep=dataset['charges']
```

```
[6]: #split into training set and test
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(indep, dep, test_size = 1/
↳3, random_state = 0)
```

```
[7]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
[8]: from sklearn.ensemble import RandomForestRegressor
```

```
[10]: from sklearn.model_selection import GridSearchCV
#from sklearn.tree import DecisionTreeRegressor
param_grid = {'criterion':['mse','mae'],
              'max_features': ['auto','sqrt','log2'],
              'n_estimators':[10,100]}

grid = GridSearchCV(RandomForestRegressor(), param_grid, refit = True, verbose=
↳ 3,n_jobs=-1)

# fitting the model for grid search
grid.fit(X_train, y_train)
```

Fitting 3 folds for each of 12 candidates, totalling 36 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done   8 tasks      | elapsed:    4.9s
[Parallel(n_jobs=-1)]: Done  26 out of  36 | elapsed:    5.5s remaining:   2.1s
[Parallel(n_jobs=-1)]: Done  36 out of  36 | elapsed:    7.0s finished
```

```
[10]: GridSearchCV(cv='warn', error_score='raise-deprecating',
                  estimator=RandomForestRegressor(bootstrap=True, criterion='mse',
                                                    max_depth=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators='warn', n_jobs=None,
                                                    oob_score=False, random_state=None,
                                                    verbose=0, warm_start=False),
                  iid='warn', n_jobs=-1,
```

```

param_grid={'criterion': ['mse', 'mae'],
            'max_features': ['auto', 'sqrt', 'log2'],
            'n_estimators': [10, 100]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=None, verbose=3)

```

```

[11]: # print best parameter after tuning
      #print(grid.best_params_)
      re=grid.cv_results_
      #print(re)
      grid_predictions = grid.predict(X_test)

      # print classification report
      from sklearn.metrics import r2_score
      r_score=r2_score(y_test,grid_predictions)

      print("The R_score value for best parameter {}: ".format(grid.
      ↪best_params_),r_score)

```

The R\_score value for best parameter {'criterion': 'mae', 'max\_features': 'sqrt', 'n\_estimators': 100}: 0.87438313765281

```

[12]: table=pd.DataFrame.from_dict(re)

```

```

[13]: table

```

```

[13]:    mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
0      0.064955    0.004217      0.004987      1.946680e-07
1      0.504170    0.003438      0.025830      1.421202e-03
2      0.049131    0.005116      0.005000      8.146892e-04
3      0.443521    0.018677      0.022813      1.652083e-03
4      0.048582    0.002549      0.005331      9.494264e-04
5      0.437335    0.019040      0.022984      1.642447e-03
6      0.221658    0.013660      0.004001      1.596079e-05
7      1.746727    0.027610      0.024313      4.391143e-04
8      0.124352    0.011595      0.003989      2.247832e-07
9      1.127960    0.011441      0.022014      8.171223e-04
10     0.120176    0.005492      0.004501      4.195174e-04
11     1.134786    0.008669      0.022513      1.459310e-03

```

```

      param_criterion param_max_features param_n_estimators  \
0                mse                auto                10
1                mse                auto                100
2                mse                sqrt                10
3                mse                sqrt                100
4                mse                log2                10
5                mse                log2                100

```

6	mae	auto	10
7	mae	auto	100
8	mae	sqrt	10
9	mae	sqrt	100
10	mae	log2	10
11	mae	log2	100

	params	split0_test_score \
0	{'criterion': 'mse', 'max_features': 'auto', '...	0.772643
1	{'criterion': 'mse', 'max_features': 'auto', '...	0.790400
2	{'criterion': 'mse', 'max_features': 'sqrt', '...	0.776431
3	{'criterion': 'mse', 'max_features': 'sqrt', '...	0.806334
4	{'criterion': 'mse', 'max_features': 'log2', '...	0.785895
5	{'criterion': 'mse', 'max_features': 'log2', '...	0.805157
6	{'criterion': 'mae', 'max_features': 'auto', '...	0.736597
7	{'criterion': 'mae', 'max_features': 'auto', '...	0.791955
8	{'criterion': 'mae', 'max_features': 'sqrt', '...	0.787623
9	{'criterion': 'mae', 'max_features': 'sqrt', '...	0.808999
10	{'criterion': 'mae', 'max_features': 'log2', '...	0.765594
11	{'criterion': 'mae', 'max_features': 'log2', '...	0.800216

	split1_test_score	split2_test_score	mean_test_score	std_test_score \
0	0.765234	0.790219	0.776028	0.010473
1	0.795172	0.798572	0.794710	0.003353
2	0.808969	0.761584	0.782321	0.019779
3	0.803550	0.791685	0.800529	0.006352
4	0.786061	0.780230	0.784064	0.002710
5	0.800521	0.794355	0.800017	0.004425
6	0.762197	0.771544	0.756757	0.014780
7	0.789747	0.786850	0.789520	0.002091
8	0.787097	0.774438	0.783058	0.006094
9	0.802397	0.796112	0.802510	0.005263
10	0.767119	0.765989	0.766233	0.000646
11	0.807552	0.791503	0.799758	0.006556

	rank_test_score
0	10
1	5
2	9
3	2
4	7
5	3
6	12
7	6
8	8
9	1
10	11

```
[ ]: age_input=float(input("Age:"))
      bmi_input=float(input("BMI:"))
      children_input=float(input("Children:"))
      sex_male_input=int(input("Sex Male 0 or 1:"))
      smoker_yes_input=int(input("Smoker Yes 0 or 1:"))

[ ]: Future_Prediction=grid.
      ↪predict([[age_input,bmi_input,children_input,sex_male_input,smoker_yes_input]])#
      ↪change the paramter,play with it.
      print("Future_Prediction={}".format(Future_Prediction))

[ ]:

[ ]:

[ ]:

[ ]:

[ ]:

[ ]:
```